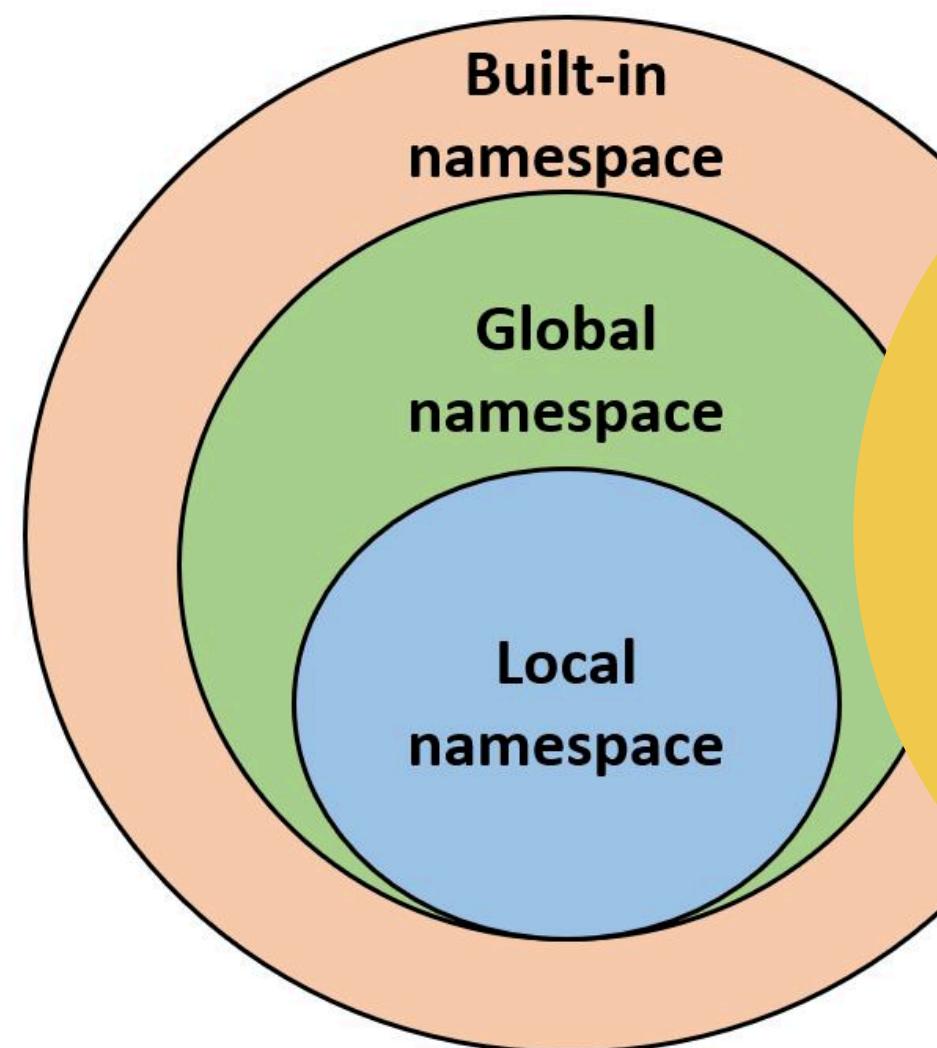
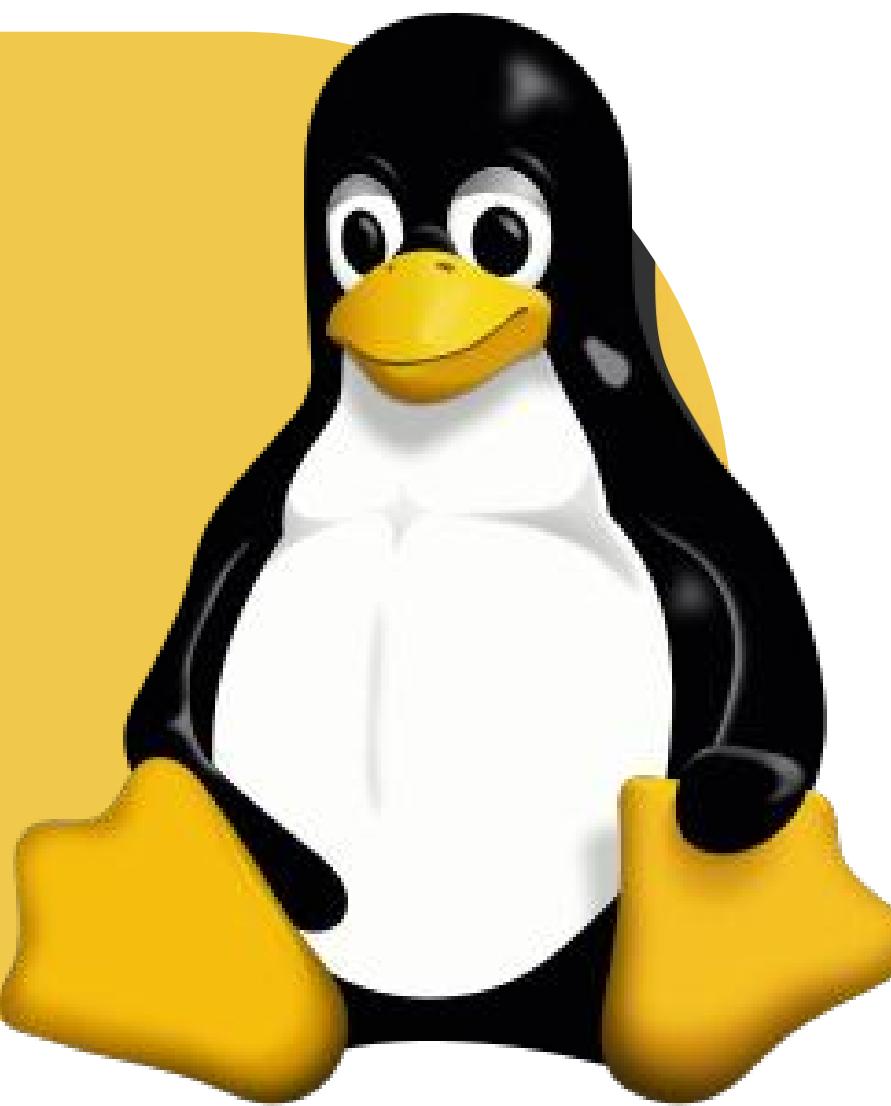


cgroups и namespaces в Linux



Cgroups



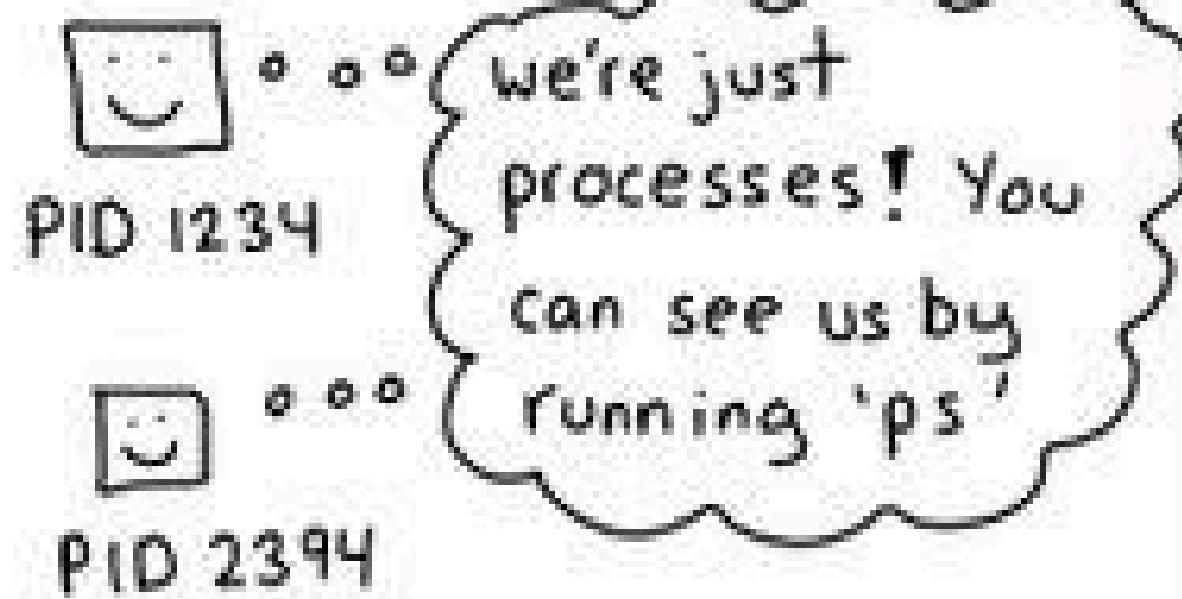
Жахонгир Ахмадалиев



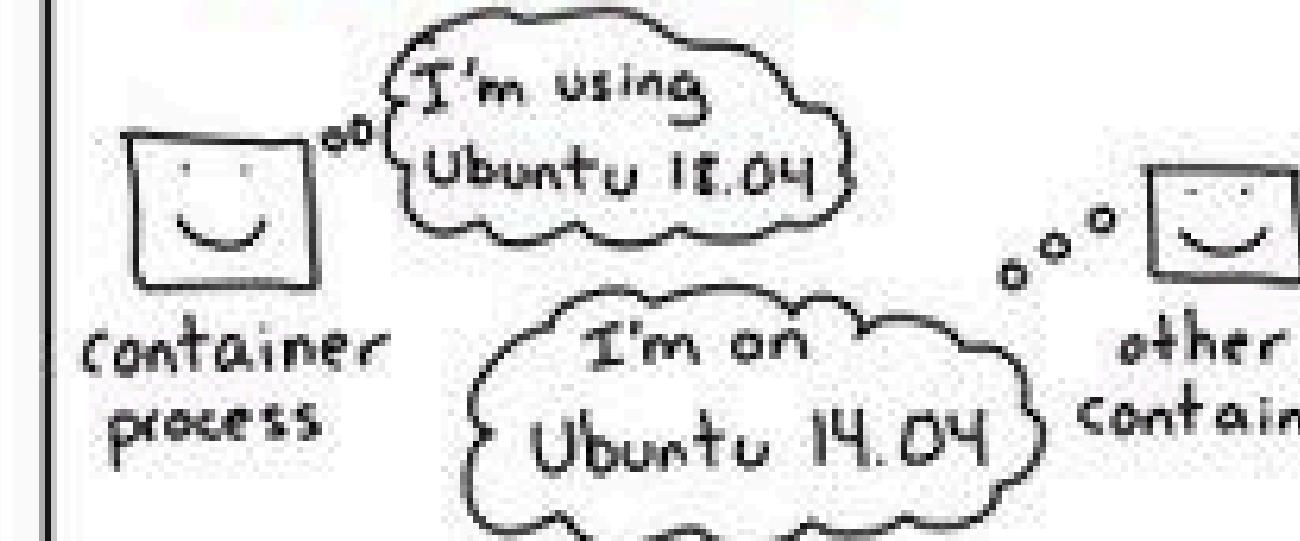
JULIA EVANS
@b0rk

namespaces

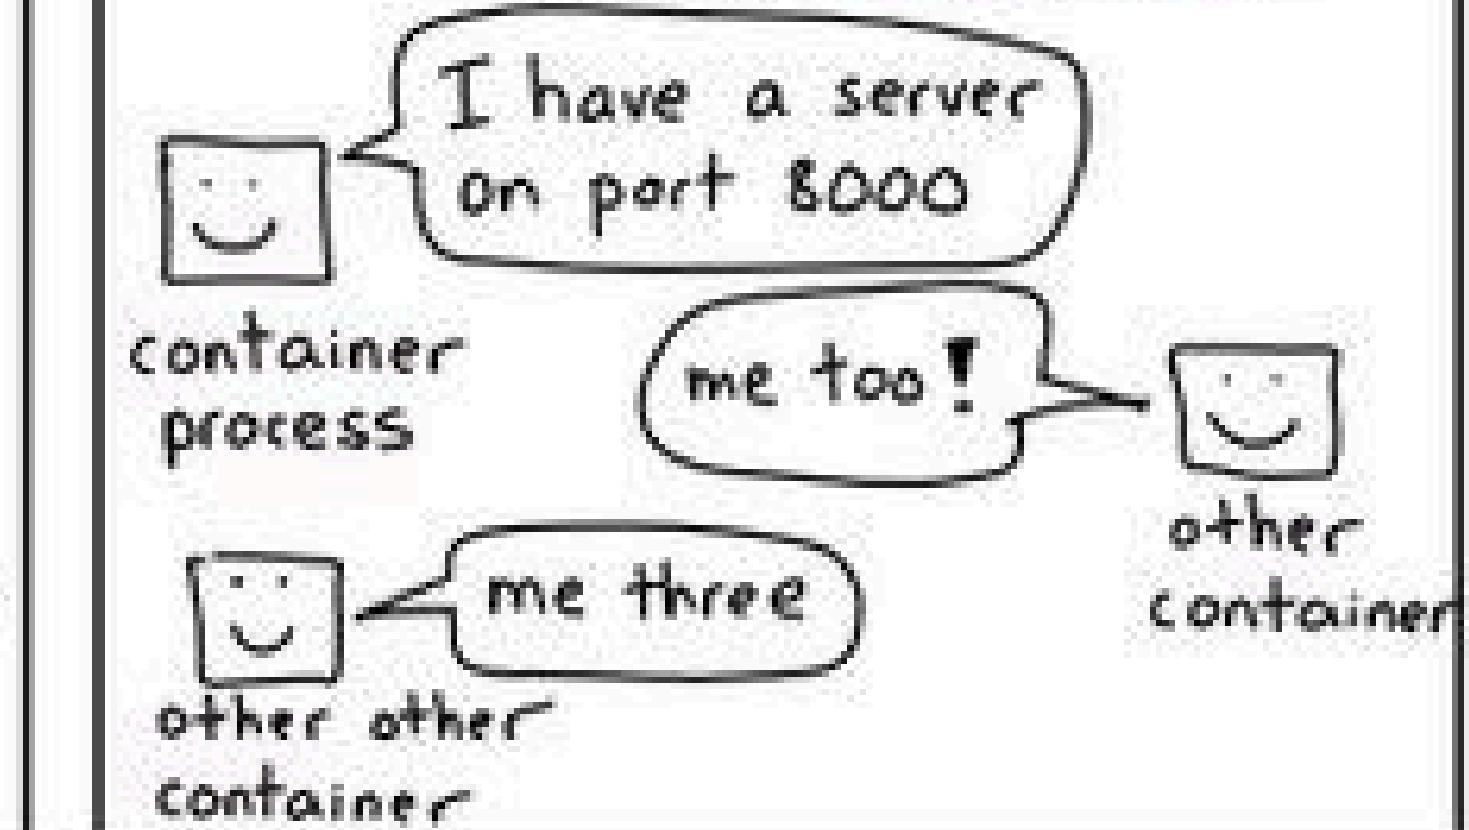
a container is a group of processes



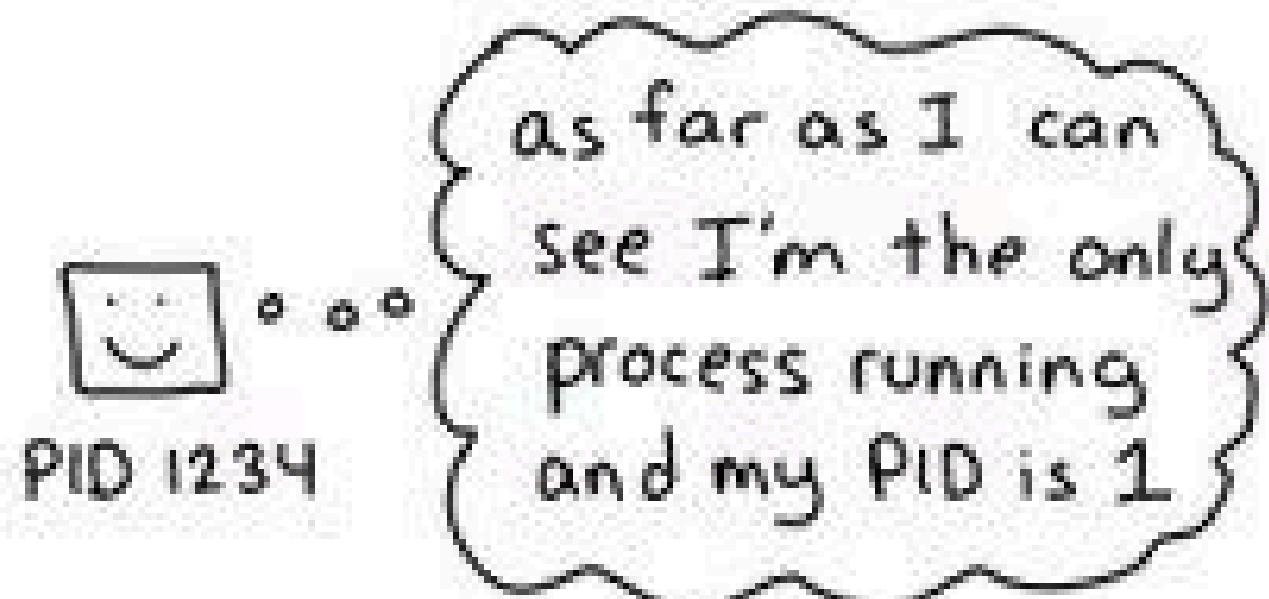
container processes have their own filesystem



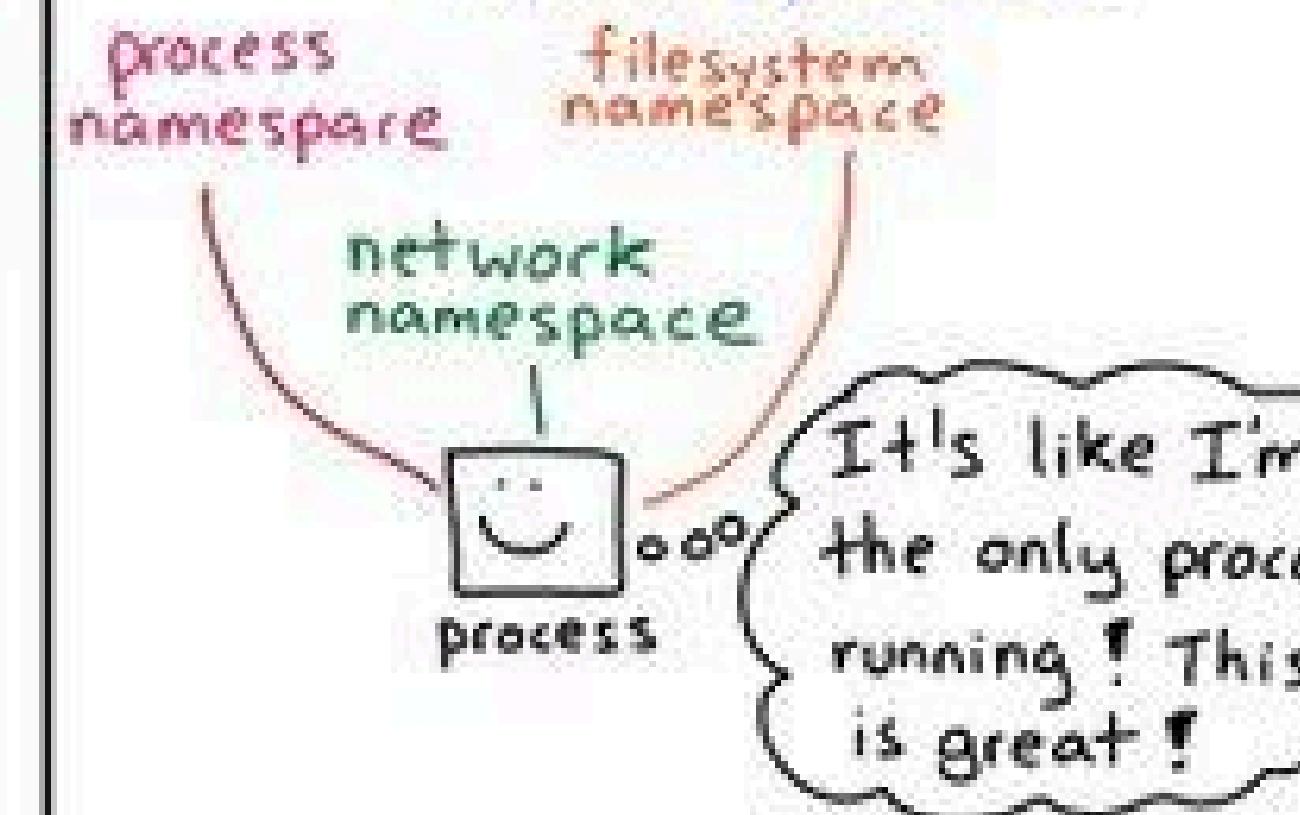
... and their own network



... and their own process list



these are called namespaces

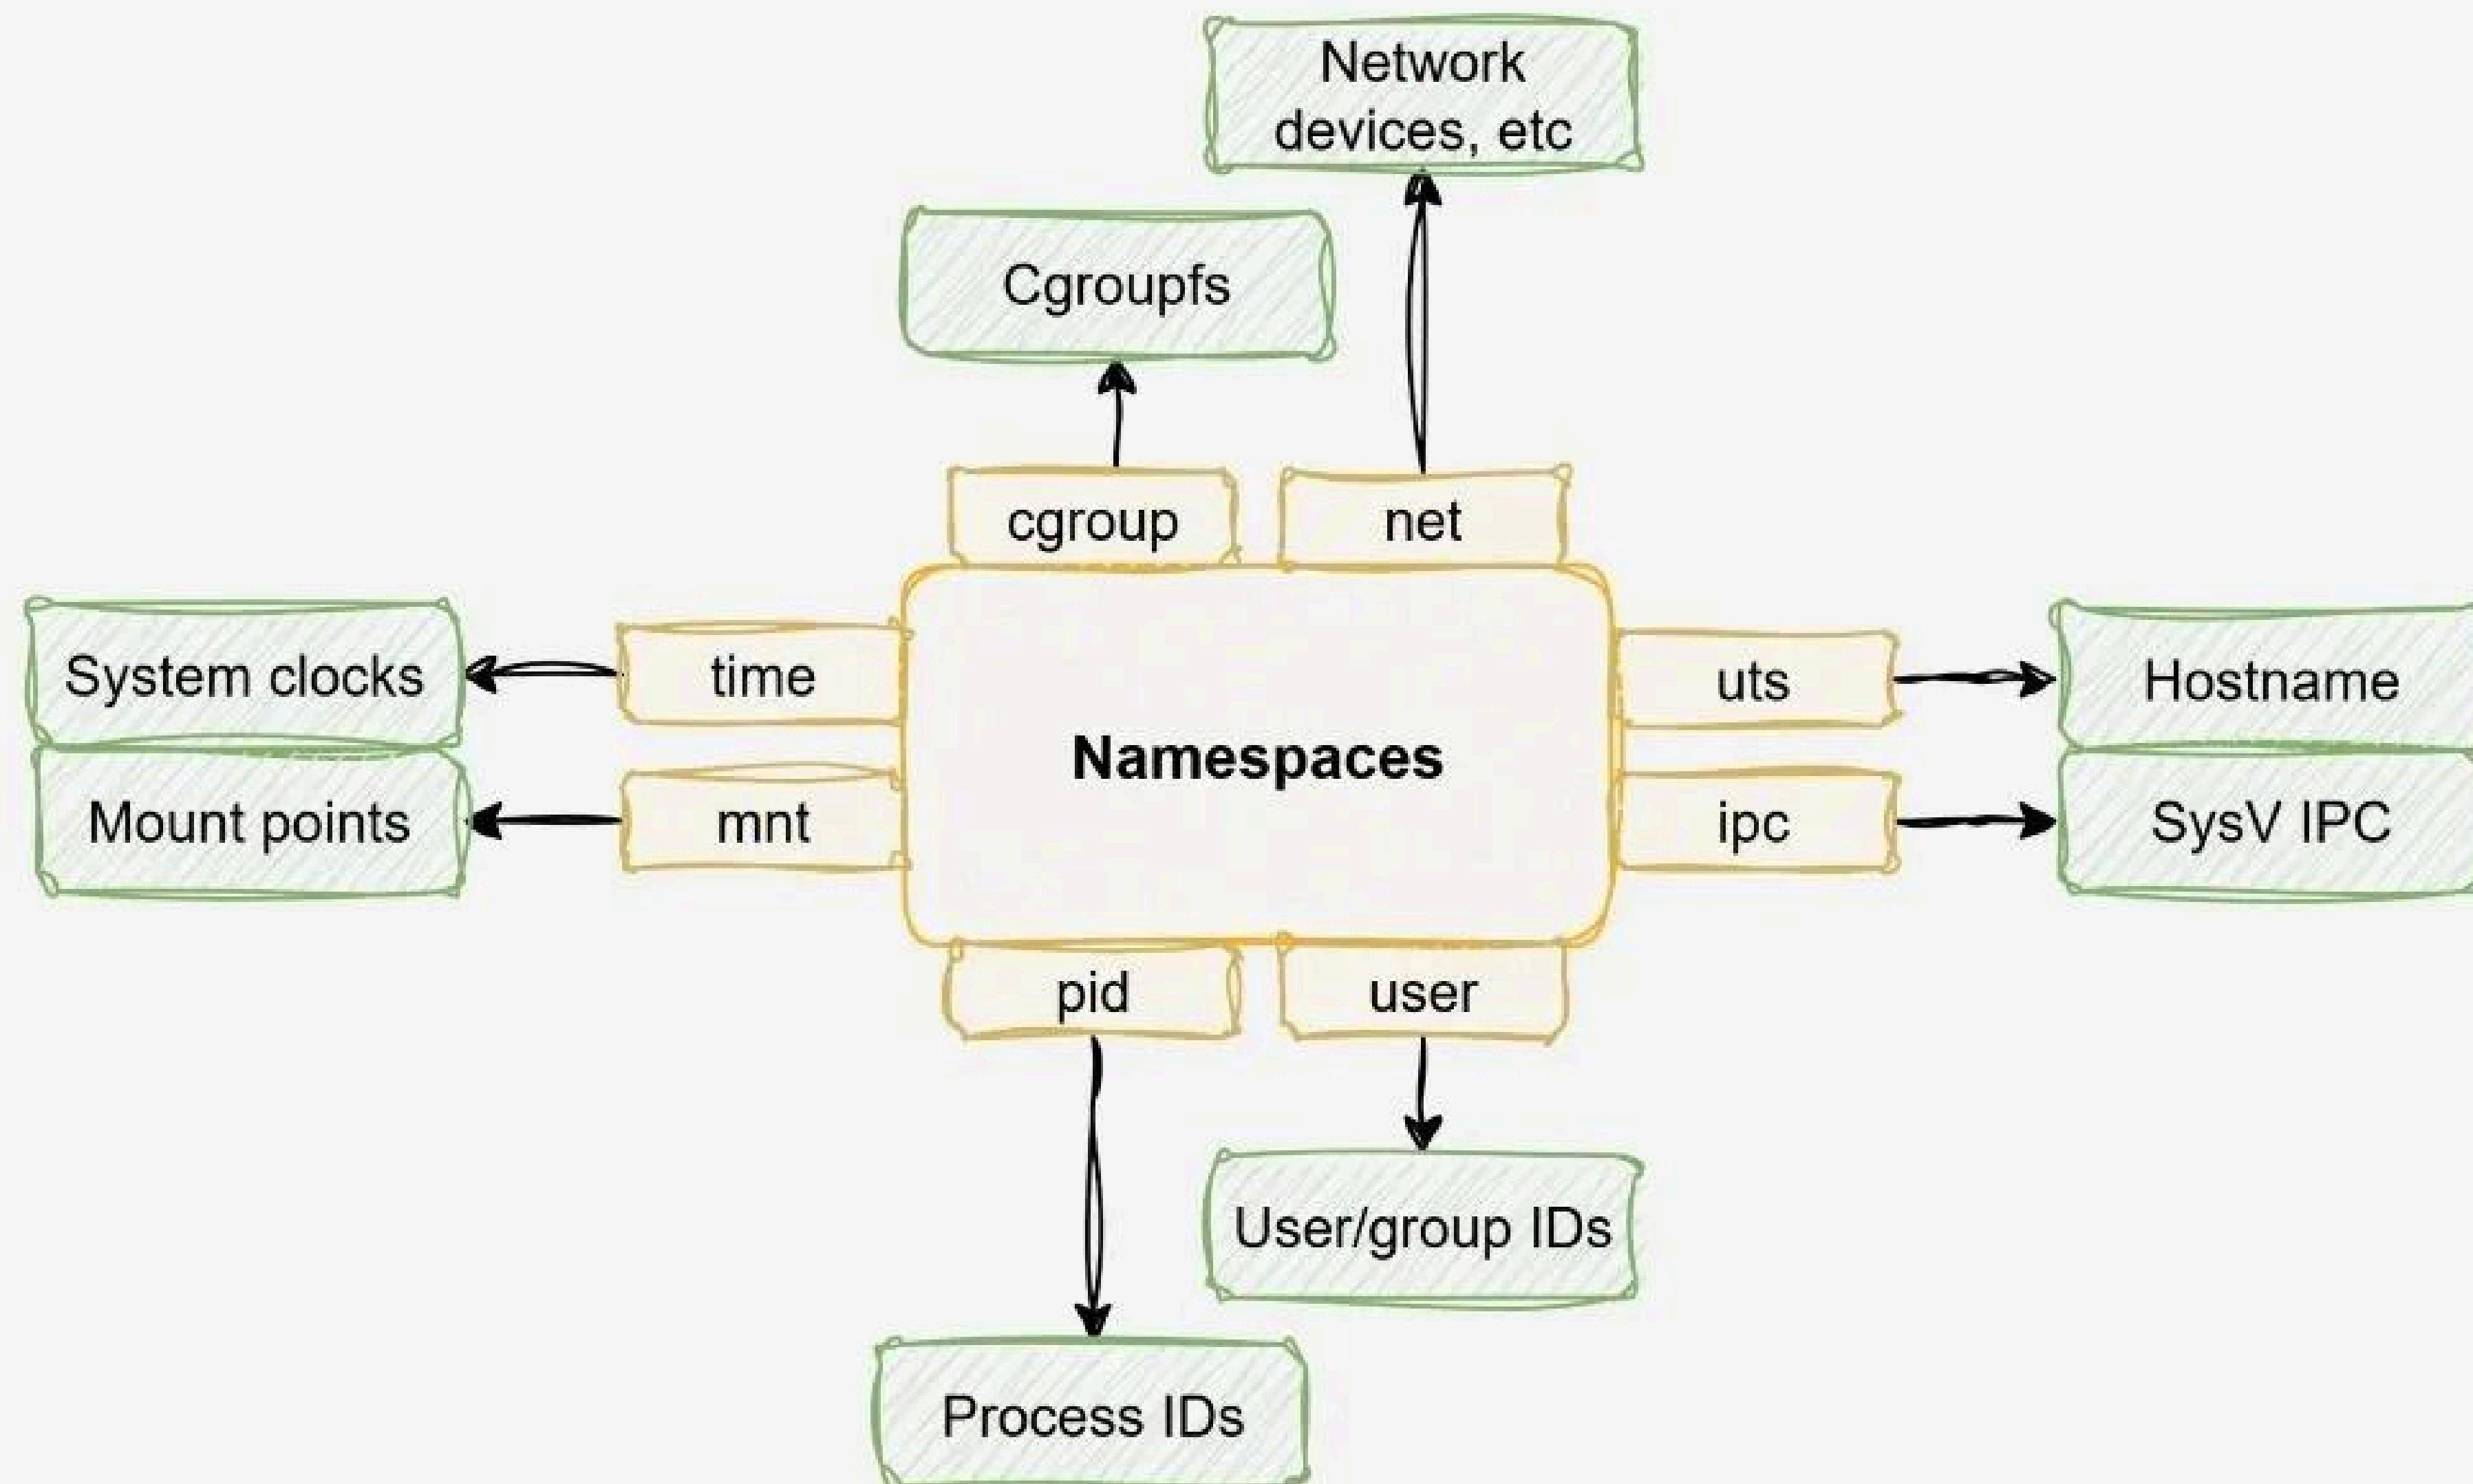


nsenter: spy on a process's namespace

This shows what ports are open in PID 1234's container:

```
$ sudo nsenter -t 1234 -n netstat - tulp
```

use PID 1234's namespaces
+ network namespace
netstat - tulp
+ command to run

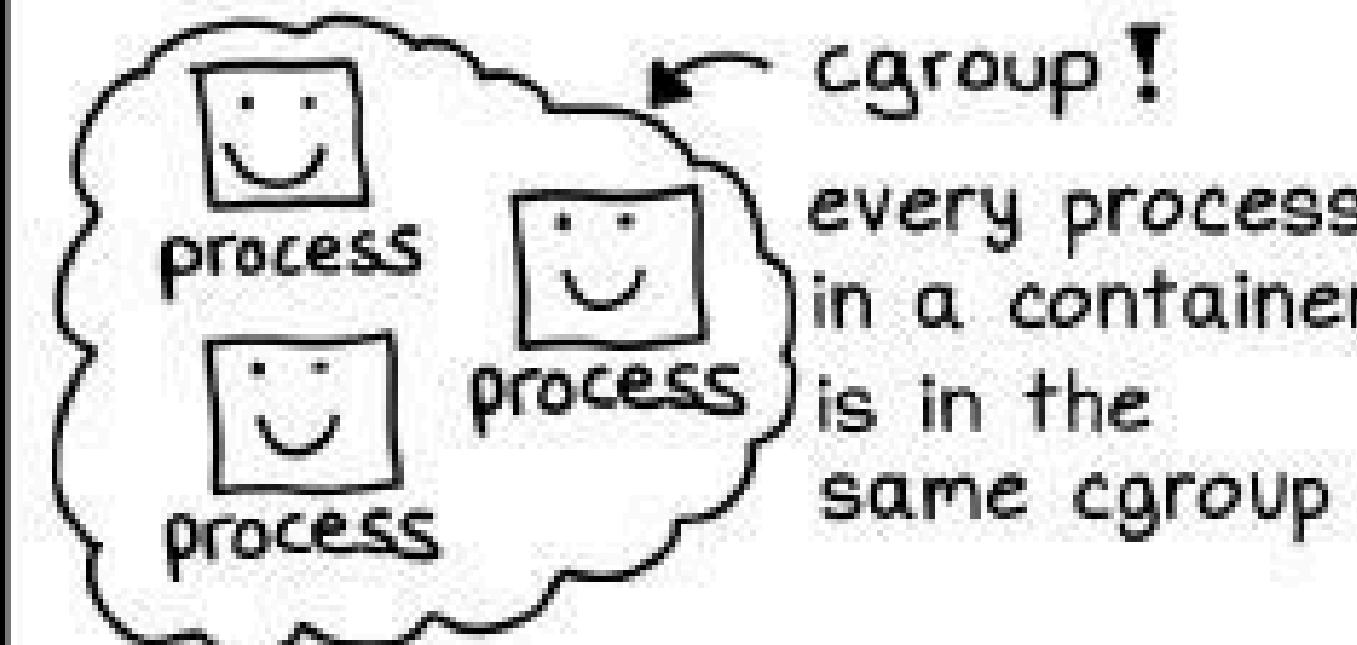


cgroups

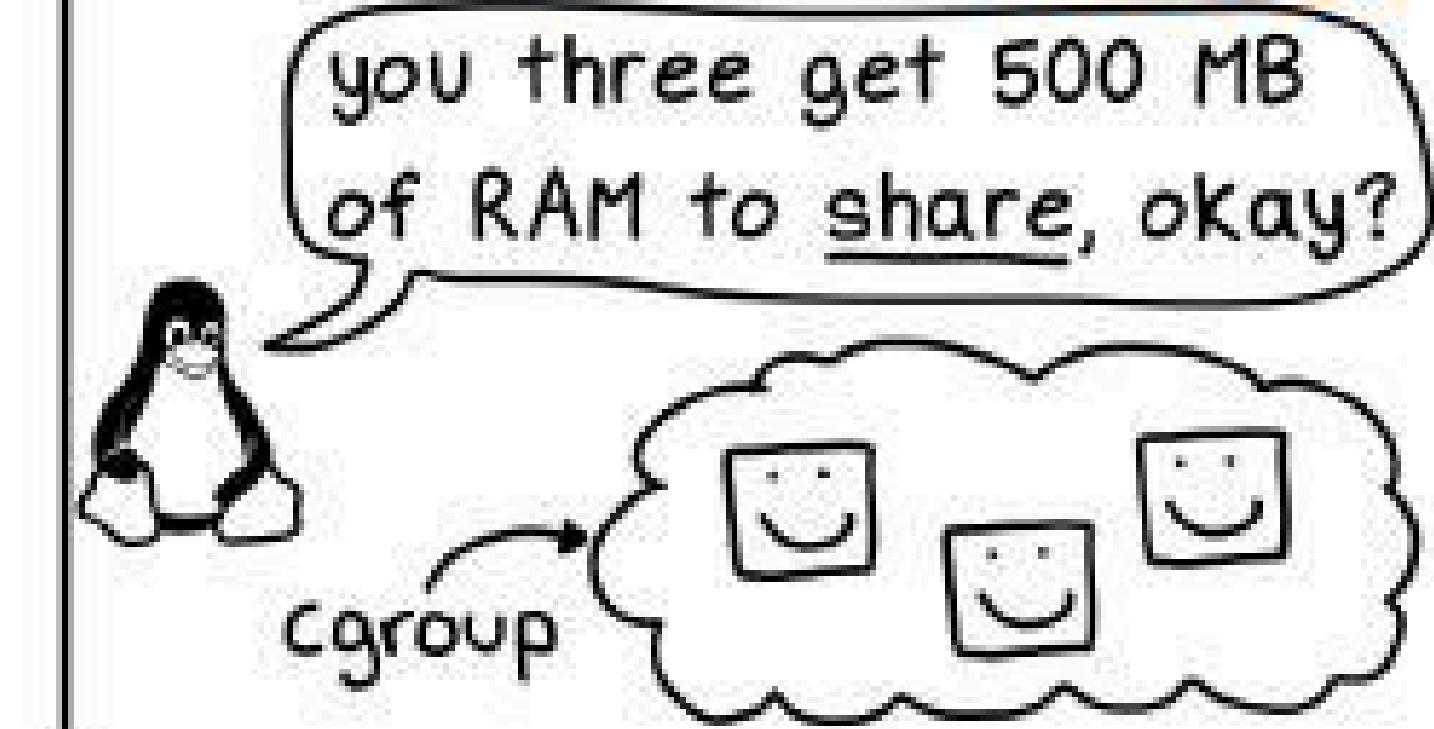
processes can use a lot of memory



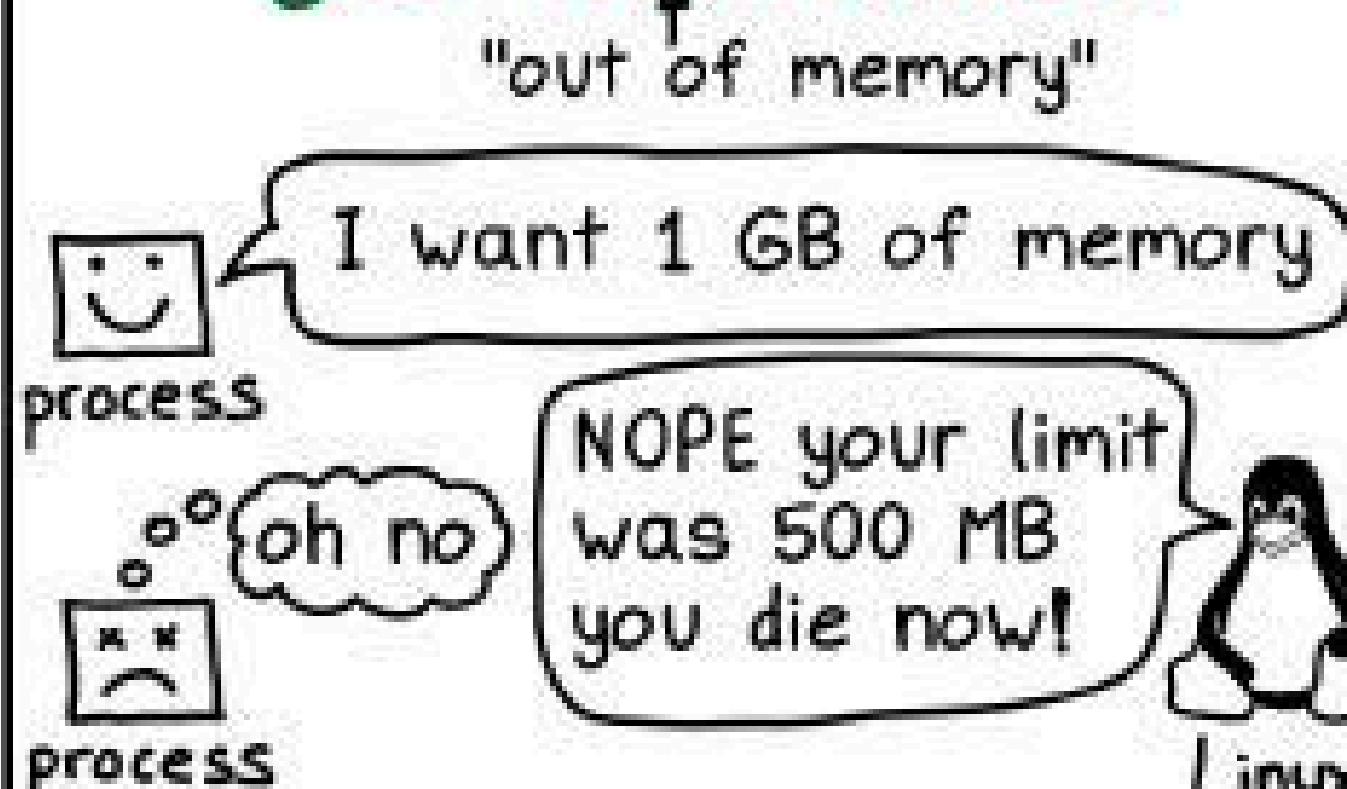
a cgroup is a group of processes



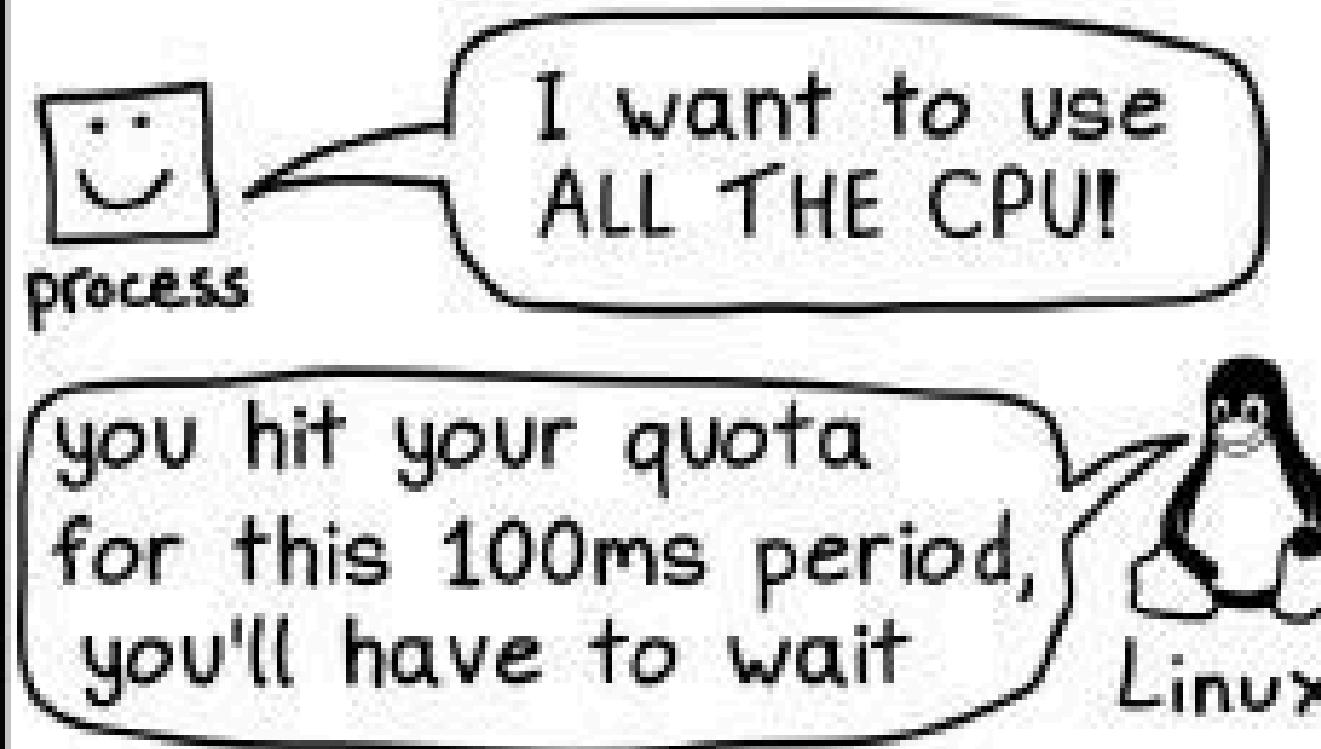
cgroups have memory/CPU limits



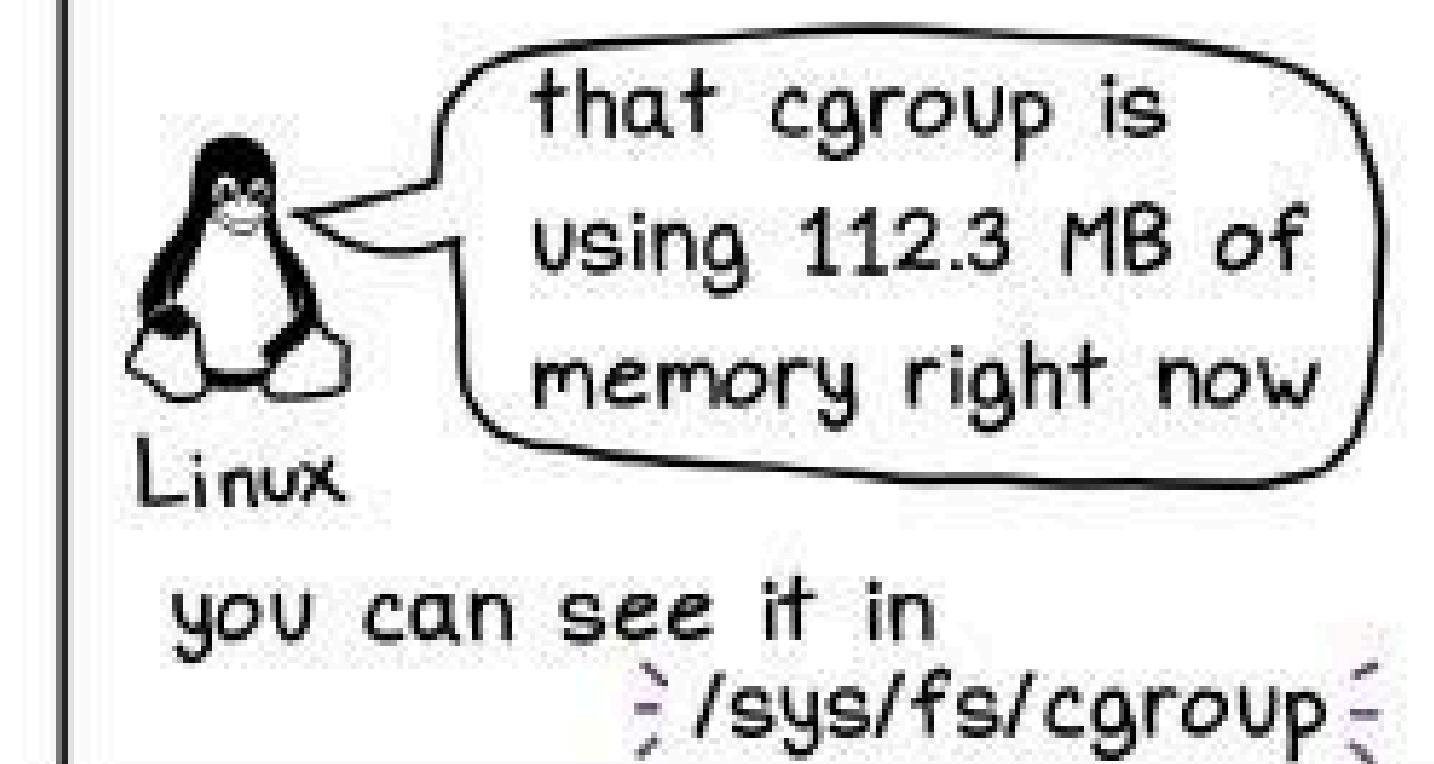
use too much memory: get OOM killed

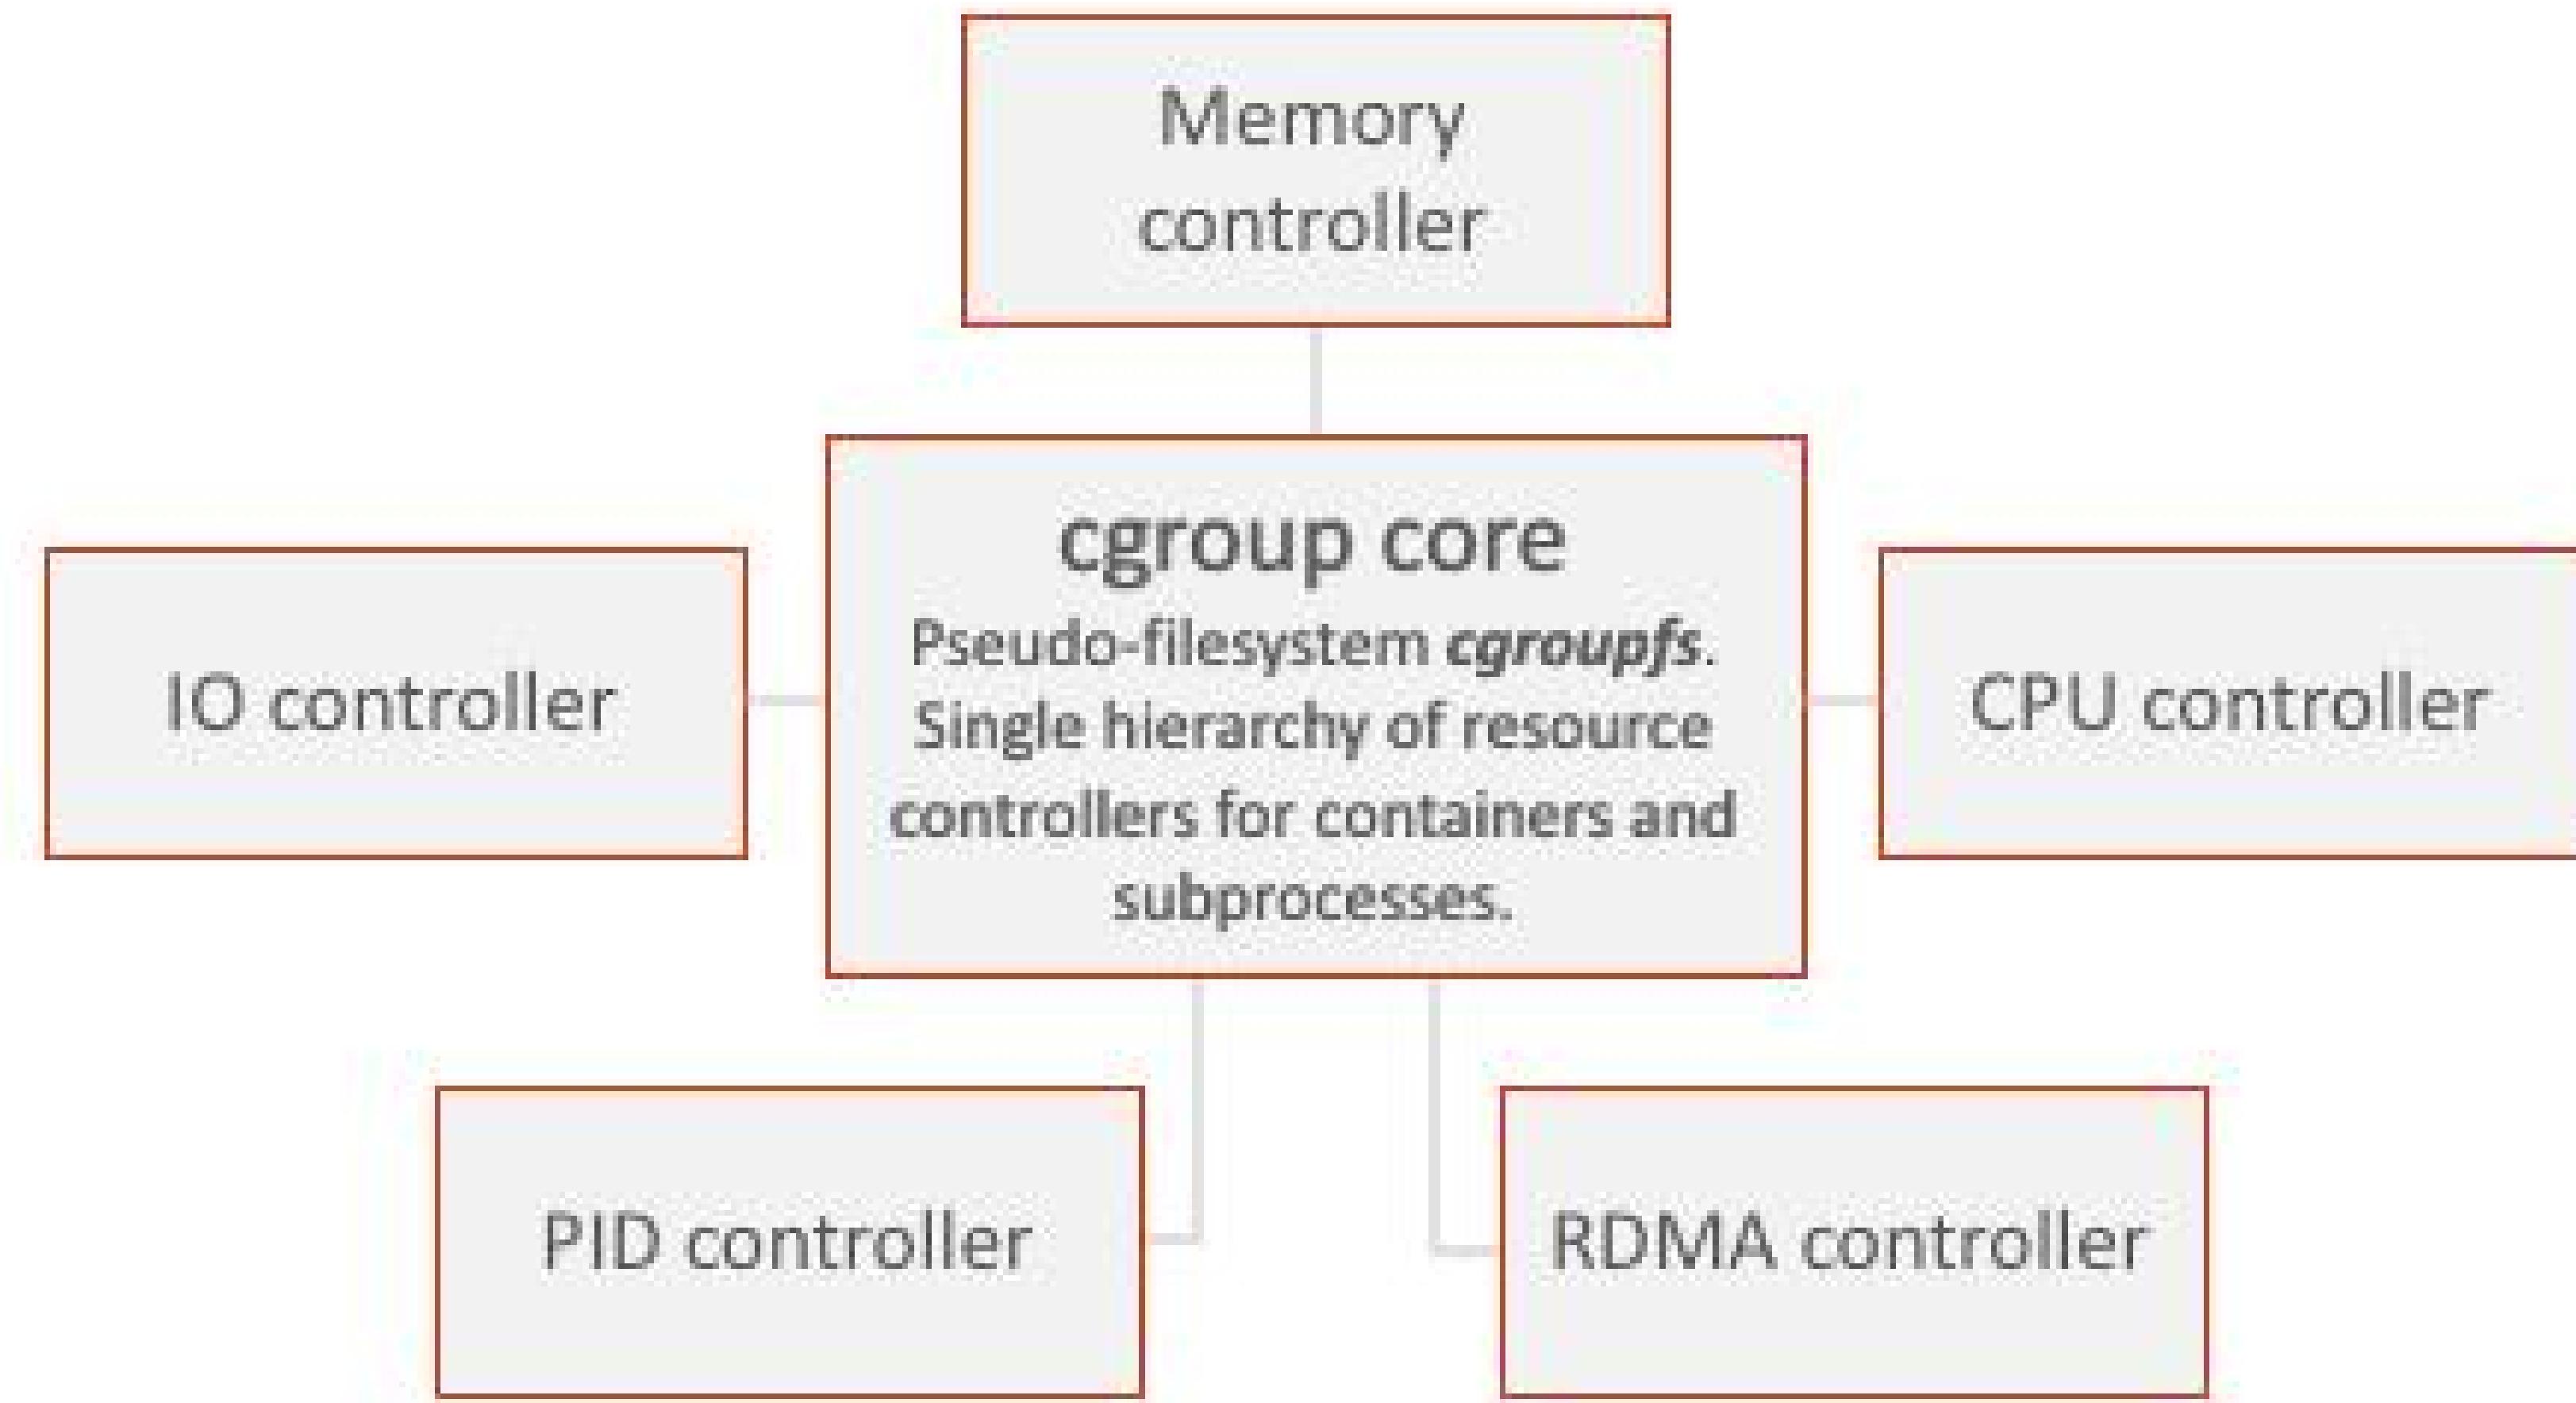


use too much CPU: get slowed down



cgroups track memory & CPU usage





NAME[top](#)

unshare - run program in new namespaces

SYNOPSIS[top](#)

unshare [options] [*program* [*arguments*]]

DESCRIPTION[top](#)

The **unshare** command creates new namespaces (as specified by the command-line options described below) and then executes the specified *program*. If *program* is not given, then "\${SHELL}" is run (default: */bin/sh*).

- `--fork` : форкает процесс после создания пространств имён.
- `--pid` : создаёт новый PID namespace.
- `--mount` : создаёт новый Mount namespace.
- `--uts` : создаёт новый UTS namespace.
- `--ipc` : создаёт новый IPC namespace.
- `--net` : создаёт новый Network namespace.
- `--user` : создаёт новый User namespace.
- `--map-root-user` : маппинг root пользователя внутри пространства имён.
- `--mount-proc` : монтирует новую файловую систему `/proc` внутри нового Mount namespace.

пример создание pid namespace

```
sudo unshare --fork --pid --mount-proc /bin/bash

[root@archlinux mars]# ps aux
USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND
root        1  0.0  0.0  7968  6404 pts/1      S    13:32   0:00 /bin/bash
root        2  0.0  0.0  9852  5976 pts/1      R+   13:32   0:00 ps aux
[root@archlinux mars]#
```

network namespace

```
sudo unshare --net /bin/bash

[sudo] password for mars:
[root@archlinux mars]# ip link
1: lo: <LOOPBACK> mtu 65536 qdisc noop state DOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
[root@archlinux mars]#
```

Комбинированный пример

```
^ mars ~ ✘ v3.13.7 ◁ ◁ ◁ 13:45
→ sudo unshare --fork --pid --net --uts --ipc --mount-proc /bin/bash
[root@archlinux mars]# ps aux
USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND
root        1  0.0  0.0  7968  6380 pts/2      S  13:45   0:00 /bin/bash
root        2  0.0  0.0  9852  6004 pts/2      R+ 13:46   0:00 ps aux
[root@archlinux mars]# ip link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noop state DOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
[root@archlinux mars]# sleep 1000 &
[1] 4
[root@archlinux mars]# ps aux
USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND
root        1  0.0  0.0  7968  6384 pts/2      S  13:45   0:00 /bin/bash
root        4  0.0  0.0  5744  3916 pts/2      S  13:47   0:00 sleep 1000
root        5  0.0  0.0  9852  6004 pts/2      R+ 13:48   0:00 ps aux
[root@archlinux mars]#
```

```
^ mars ~ ✘ v3.13.7 ◁ ◁ ◁ 13:47
→ ps aux | grep zsh
mars      4914  0.1  0.0  20980 18724 pts/0      Ss 13:06   0:02 /usr/bin/zsh
mars      9798  1.0  0.0  20944 18688 pts/1      Ss 13:45   0:01 /usr/bin/zsh
mars     10663  0.0  0.0   6620  4160 pts/1      S+ 13:48   0:00 grep --color=auto zsh
^ mars ~ ✘ v3.13.7 ◁ ◁ ◁ 13:48
→ ip link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default
    qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: wlan0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP mode DORMANT group
    default qlen 1000
    link/ether 00:45:e2:cb:45:65 brd ff:ff:ff:ff:ff:ff
3: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN mode DEFAULT
    group default
    link/ether f2:bb:e7:71:5f:76 brd ff:ff:ff:ff:ff:ff
4: br-0e33d5dbd8db: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN mode
    DEFAULT group default
    link/ether a2:3e:5f:3d:6f:95 brd ff:ff:ff:ff:ff:ff
5: br-5b4a454eeb3c: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN mode
    DEFAULT group default
    link/ether ae:ca:74:a6:50:b0 brd ff:ff:ff:ff:ff:ff
^ mars ~ ✘ v3.13.7 ◁ ◁ ◁ 13:48
→ ps aux | grep sleep
root     10436  0.0  0.0   5744  3916 pts/2      S  13:47   0:00 sleep 1000
mars     10811  0.0  0.0   6620  4156 pts/1      S+ 13:48   0:00 grep --color=auto sleep
^ mars ~ ✘ v3.13.7 ◁ ◁ ◁ 13:48
→
```

```
▲ mars ~ ✘ v3.13.7 ☺ 15:52
```

```
→ sudo mkdir /sys/fs/cgroup/demo  
[sudo] password for mars:
```

```
▲ mars ~ ✘ v3.13.7 ☺ 15:52
```

```
→ echo "10485760" | sudo tee /sys/fs/cgroup/demo/memory.max  
10485760
```

```
▲ mars ~ ✘ v3.13.7 ☺ 15:52
```

```
→ sudo bash -c "echo \$\$ > /sys/fs/cgroup/demo/cgroup.procs && python3"
```

```
Python 3.13.7 (main, Aug 15 2025, 12:34:02) [GCC 15.2.1 20250813] on linux  
Type "help", "copyright", "credits" or "license" for more information.
```

```
>>> data = []  
>>>  
>>> while True:  
...     data.append(' ' * 1000000)  
...  
...
```

```
zsh: killed      sudo bash -c "echo \$\$ > /sys/fs/cgroup/demo/cgroup.procs && python3"
```

```
▲ mars ~ ✘ v3.13.7 ☺ 15:54
```

<https://github.com/Jahamars/pycontainer>



```
→ sudo python main.py shell
```

```
=====
```

Контейнер: demo

```
=====
```

- ✓ CGroup: 50MB RAM, 25% CPU
- ✓ RootFS: /tmp/container_demo_6ou3hx4w
- ☐ Команда: /bin/sh

```
/ # echo $$  
1  
/ # ip a  
1: lo: <LOOPBACK> mtu 65536 qdisc noop state DOWN qlen 1000  
      link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00  
/ # ls  
bin  dev  etc  lib  lib64  proc  tmp  usr  
/ # █
```

